

DEEP LEARNING WITH KERAS

CNN – FASHION MNIST

YANN LECUN'S JOB @ FACEBOOK



Write a comment... Like Comment GIF Shield

Suggested Post

 Wayfair
Sponsored · 10h

Like Page ...

Don't let these items pass you by! Save up to 70% OFF and free shipping on orders over \$49 at Wayfair!



Starting at \$293.99 Shop Now

Maitland Console Table



Starting at \$582.79 >

Flex Pro Series Adjustable Stand



The Best Tee Shirt In Your Closet
buckmason.com
Butter-soft tees that fit perfectly and hold their shape after every wash.



UPLIFT Desk
upliftdesk.com
\$495 Full Desk (not a converter). Entire desktop rises with you. Free shipping!

English (US) - 한국어 - Tiếng Việt - Bahasa Indonesia - Español +

Privacy · Terms · Advertising · Ad Choices · Cookies · More · Facebook © 2018

Suggested Post, Banner Ads

IMAGE RECOGNITION

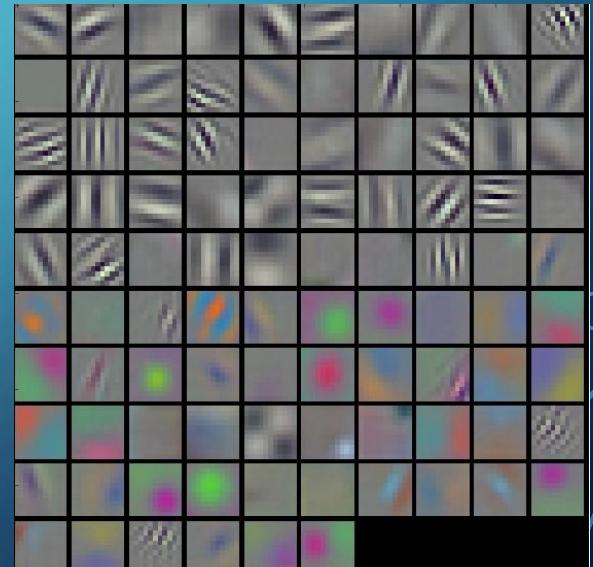
- Crucial in many ML applications



- Main challenge: # of pixels!!!!!!

CONVOLUTIONAL NEURAL NETWORK (CNN)

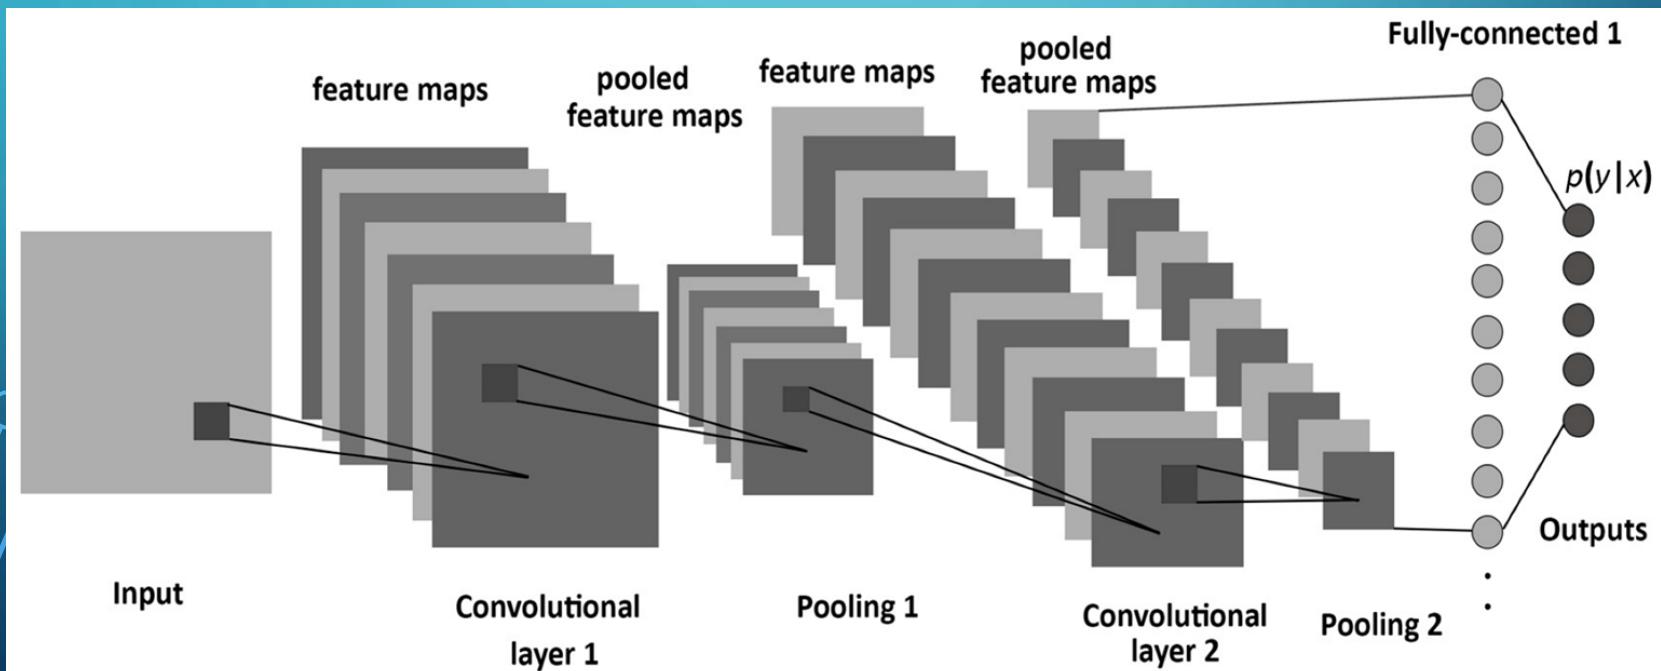
- What is Convolution?
 - 영상으로부터 특정 feature들을 추출하기 위한 필터를 구현할 때 convolution을 사용
- Why CNN?
 - Yann LeCun (1989)
 - 이미지 인식에서 MLP (Multi-Layer Perceptron)를 사용하게 되면, MLP는 모든 입력이 위치와 상관없이 동일한 수준의 중요도를 갖음
 - 이를 이용해 fully-connected neural network를 구성하게 되면 model 크기가 엄청나게 커지는 문제가 생김
 - 이에 대한 해결책으로 탄생한 것이 CNN



CNN DEEP LEARNING

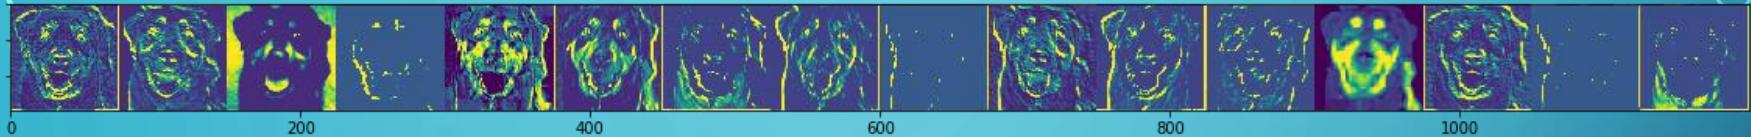
Feature Extraction

- Key idea to manage complexity while maintaining accuracy
- Features are extracted via CONV:RELU:POOL

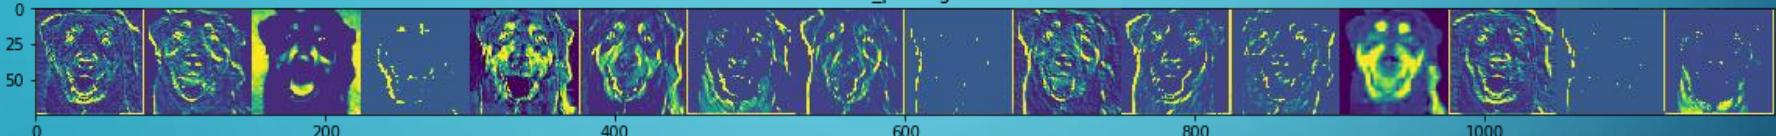


FEATURE MAP EXAMPLES

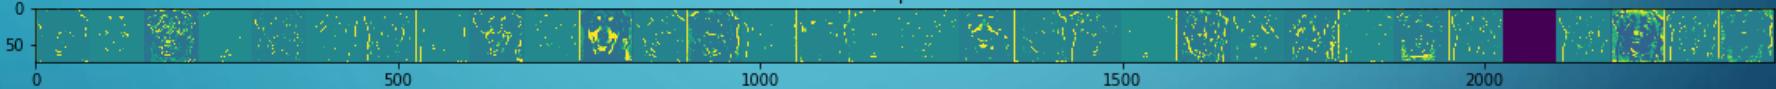
conv2d



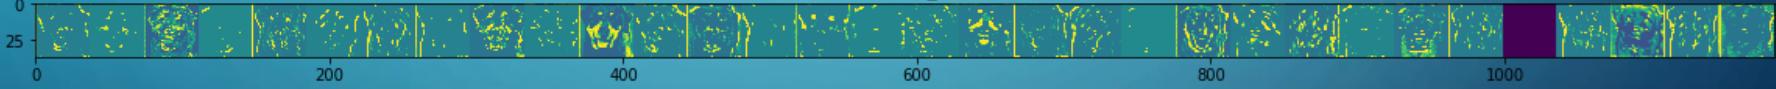
max_pooling2d



dropout



conv2d_1



max_pooling2d_1



conv2d_2

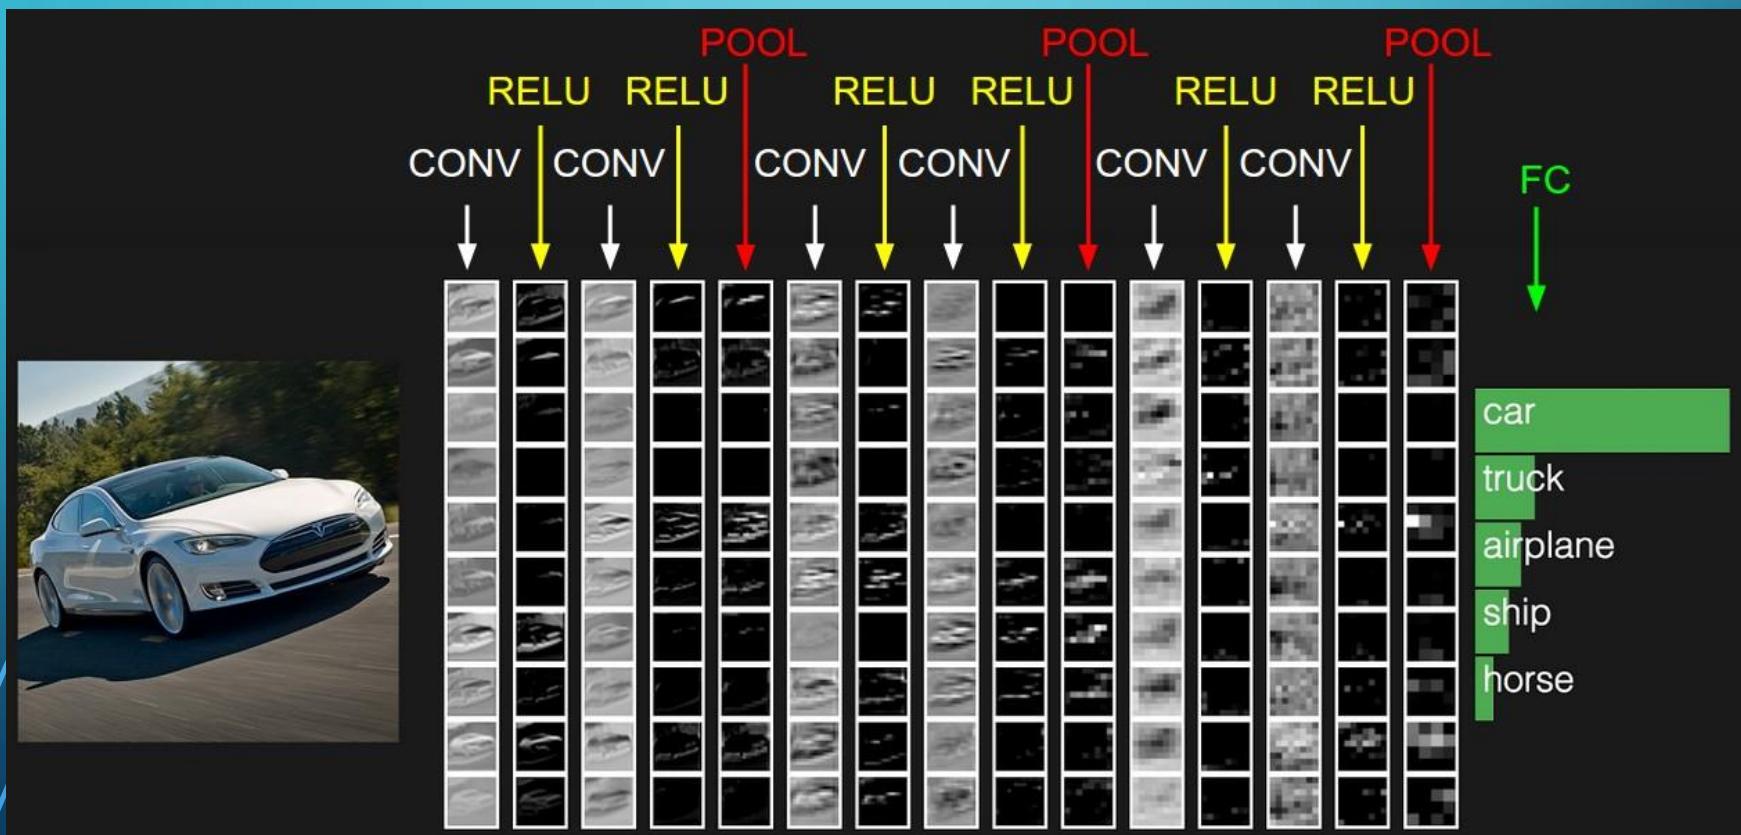


max_pooling2d_2



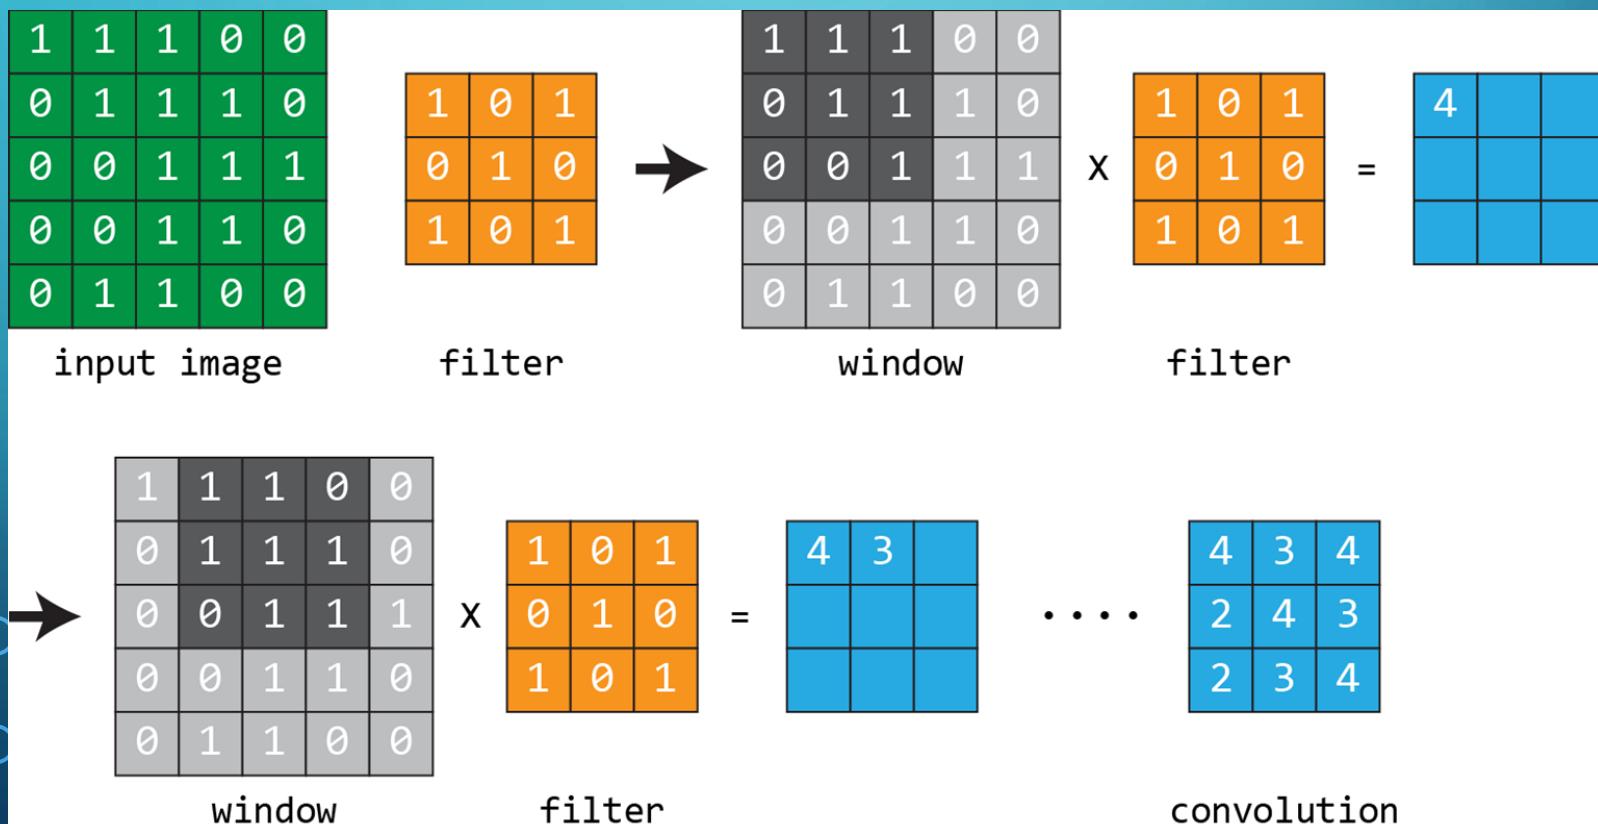
CNN DEEP LEARNING EXAMPLE

Hugely popular in image recognition



CONVOLUTION EXAMPLE

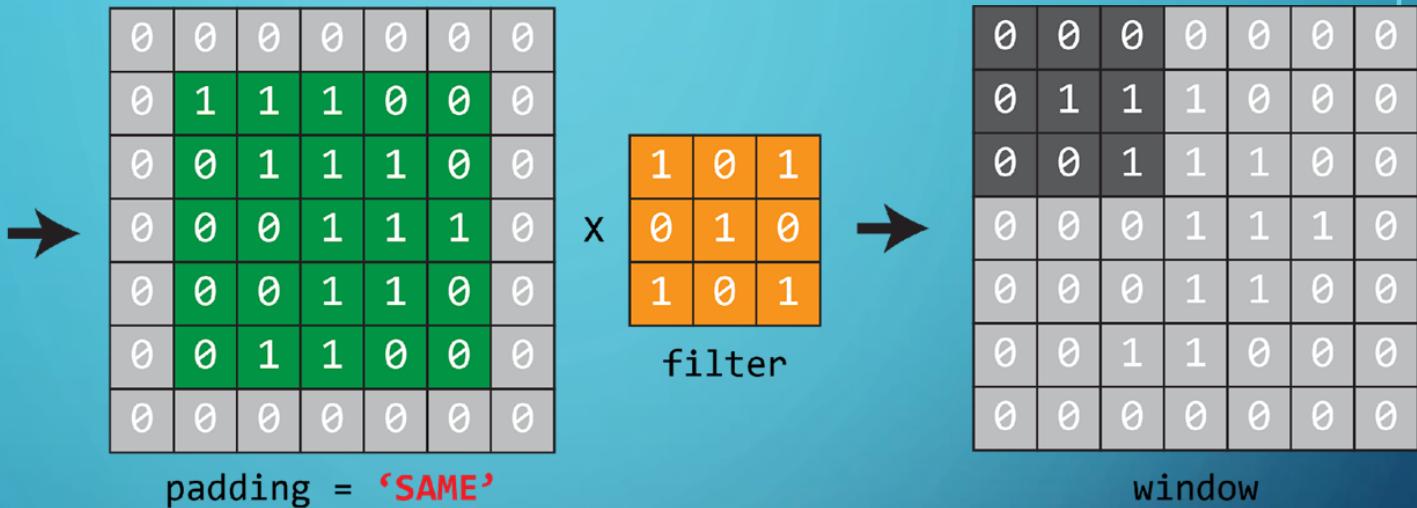
Objective: to extract features from the input data



CONVOLUTION WITH PADDING

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5x5 input image



2					

.....

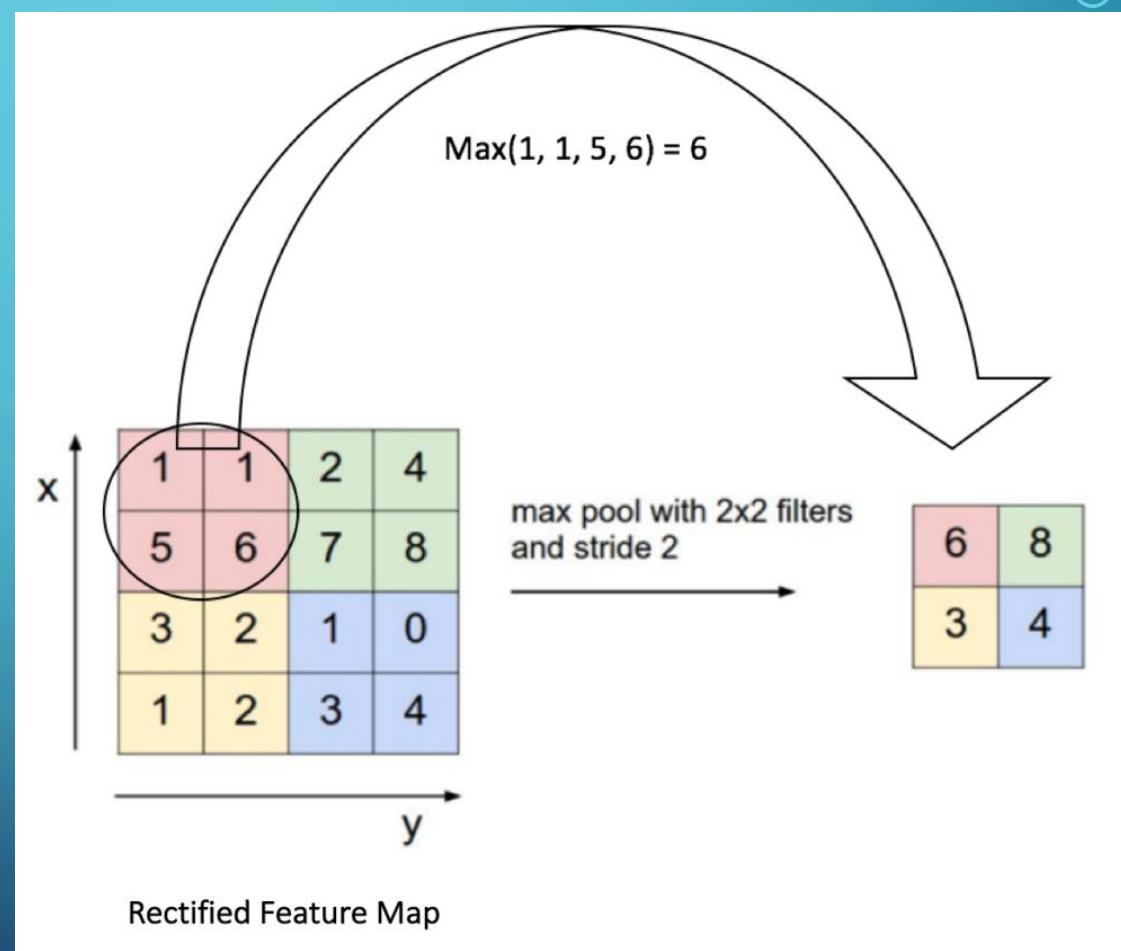
2	2	3	1	1
1	4	3	4	1
1	2	4	3	3
1	2	3	4	1
0	2	2	1	1

5x5 convolution
same size as the image

POOLING

objective

- reduce the feature map dimension and keep the important information
- types: max, average, sum, etc.



POOLING WITH PADDING

Better handle odd dimensions

4	7	3	1	6
3	3	7	9	2
6	4	5	8	9
0	1	4	3	7
2	5	7	8	3

5x5 input image
odd dimension

→ max pooling
2x2 window
stride 2
padding = ‘VALID’

4	7	3	1	6
3	3	7	9	2
6	4	5	8	9
0	1	4	3	7
2	5	7	8	3

periphery ignored

4	7	3	1	6
3	3	7	9	2
6	4	5	8	9
0	1	4	3	7
2	5	7	8	3

5x5 input image
odd dimension

→ max pooling
2x2 window
stride 2
padding = ‘SAME’

4	7	3	1	6	0
3	3	7	9	2	0
6	4	5	8	9	0
0	1	4	3	7	0
2	5	7	8	3	0
0	0	0	0	0	0

padding added

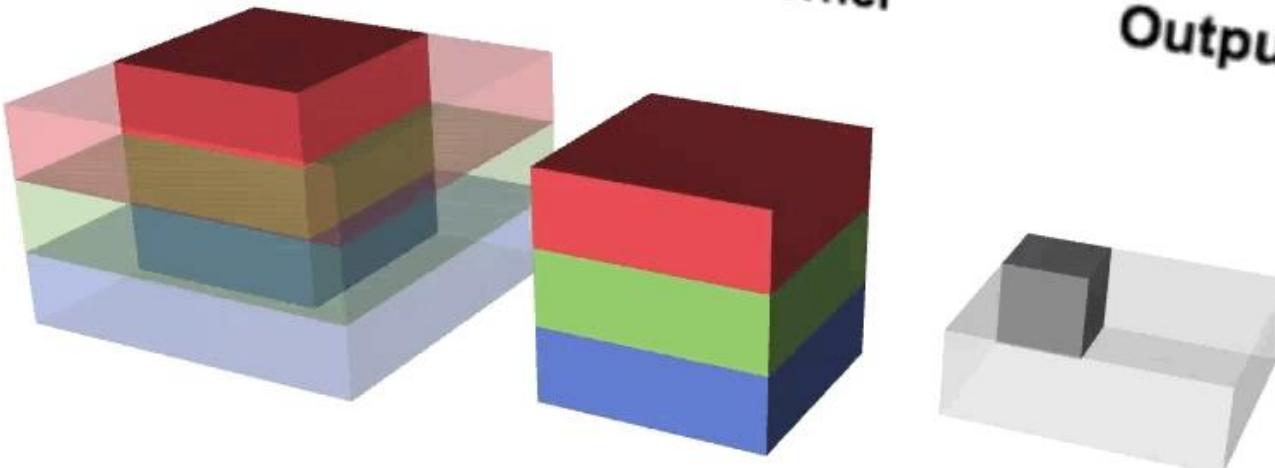
HANDLING MULTIPLE CHANNELS

3D filters

Input

Kernel

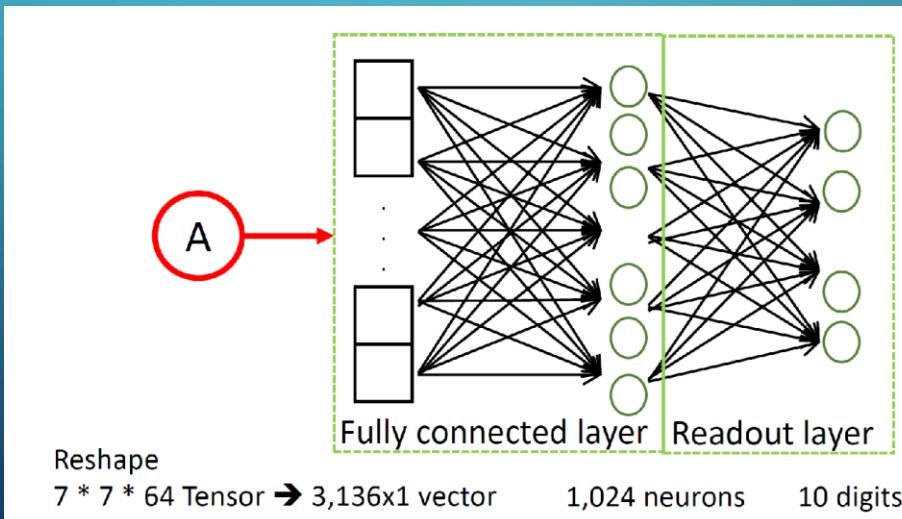
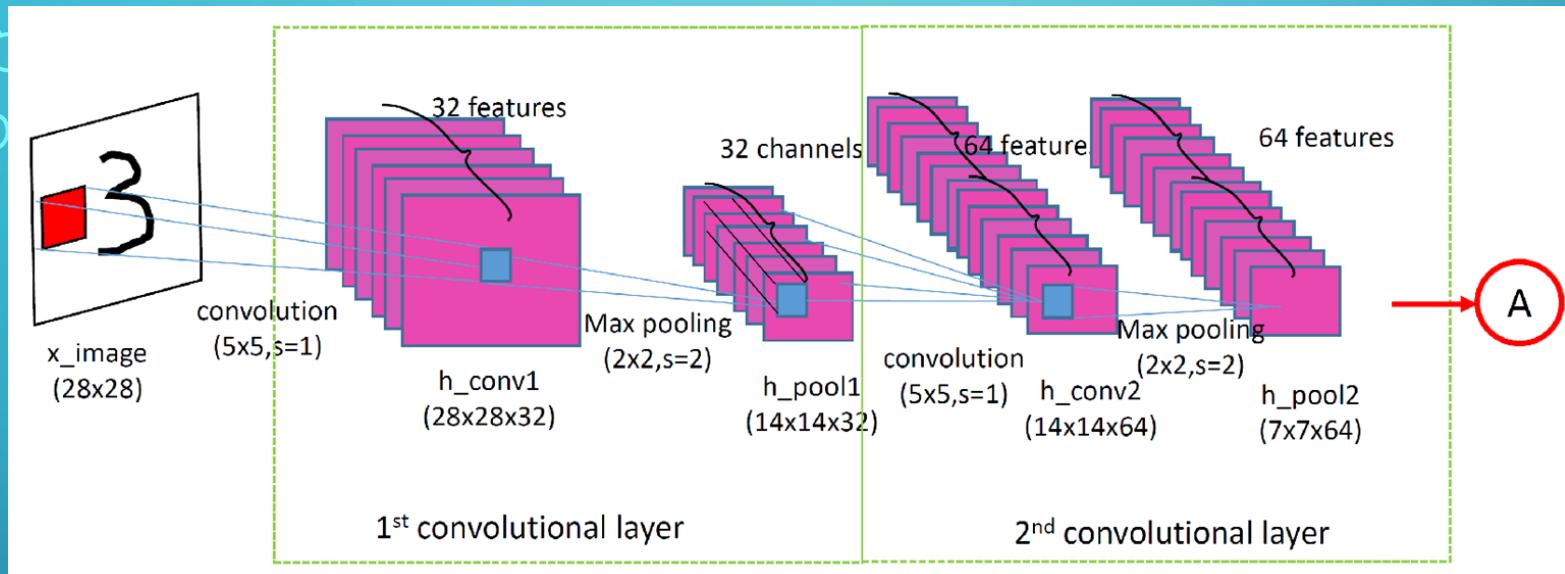
Output



<https://medium.com/apache-mxnet/multi-channel-convolutions-explained-with-ms-excel-9bbf8eb77108>

HANDLING MULTIPLE CHANNELS (CONT)

LECUN CNN ARCHITECTURE



FASHION MNIST DATABASE

10 fashion items in gray-scale images



- | | |
|--------------|--|
| 0 T-shirt | |
| 1 Trouser | |
| 2 Pullover | |
| 3 Dress | |
| 4 Coat | |
| 5 Sandal | |
| 6 Shirt | |
| 7 Sneaker | |
| 8 Bag | |
| 9 Ankle boot | |

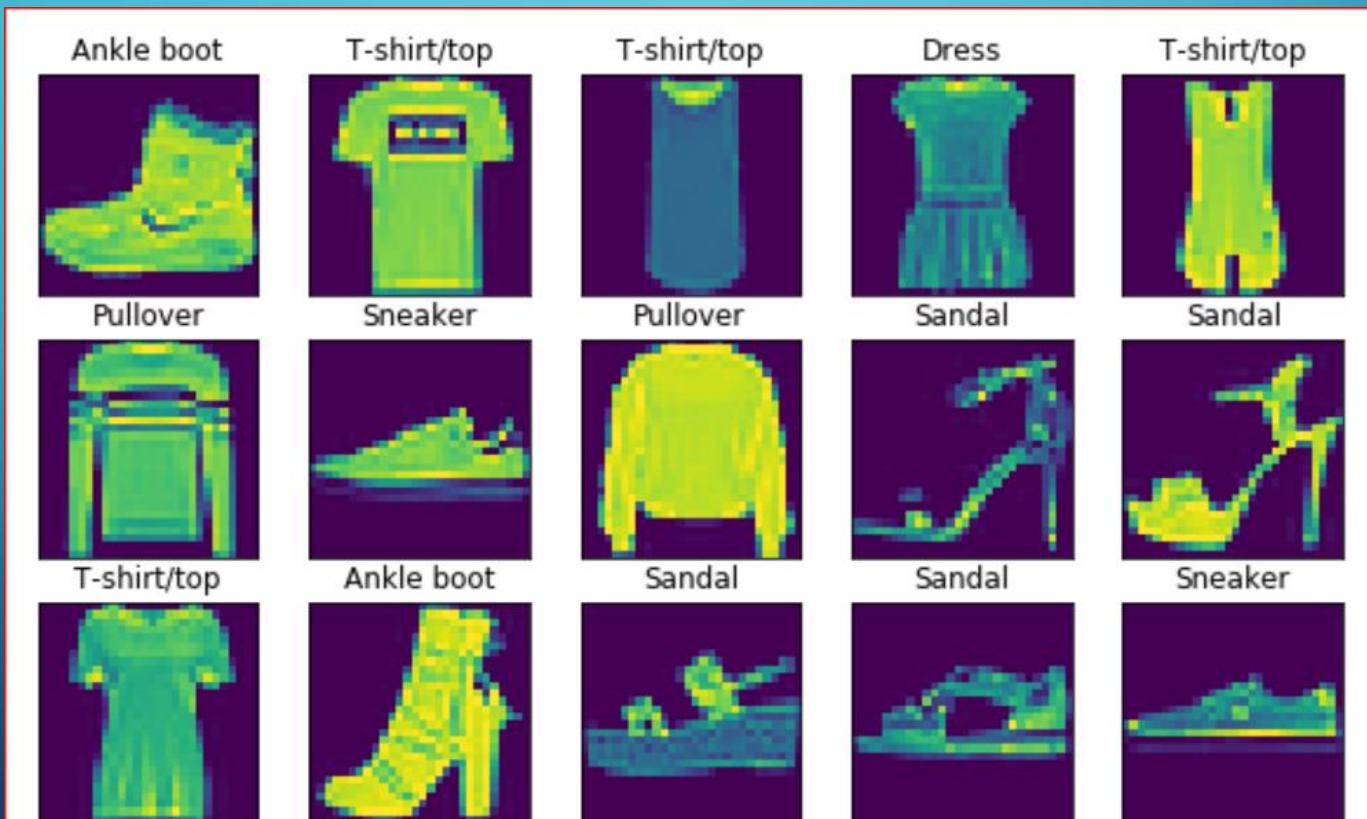
Modern
Deep Learning
Fashion-MNIST
with
Keras
TensorFlow

By Margaret 4/28/2018

FASHION MNIST DATABASE (CONT)

70,000 low resolution image files (28x28 pixels)

- Each pixel data is in [0, 255]



3. FASHION-ITEMS.PY (1/12)

```
# import packages  
import matplotlib.pyplot as plt  
import numpy as np  
from keras.datasets import fashion_mnist  
from keras.models import Sequential  
from keras.utils import np_utils  
from keras.layers import Flatten, Activation  
from keras.layers import Dense, MaxPool2D, Conv2D, InputLayer
```

3. FASHION-ITEMS.PY (2/12)

```
# global constants and hyper-parameters  
TOTAL_CLASS = 10  
MY_SAMPLE = 5  
MY_EPOCH = 10  
MY_BATCH = 200
```

3. FASHION-ITEMS.PY (3/12)

```
# load DB and split into train vs. test
# fasion MNIST data is 28x28 gray-scale image
(X_train, Y_train), (X_test, Y_test) = fashion_mnist.load_data()

print('\n== DB SHAPING INFO ==')
print("X train shape:", X_train.shape)
print("Y train shape:", Y_train.shape)
print("X test shape:", X_test.shape)
print("Y test shape:", Y_test.shape)
```

3. FASHION-ITEMS.PY (4/12)

```
# define labels
labels = ['t-shirt', 'trouser', 'pullover', 'dress', 'coat',
, 'sandal', 'shirt', 'sneaker', 'bag', 'ankle-boot']

# display a sample image and label
print('\n== SAMPLE DATA (RAW) ==')
sample = Y_train[MY_SAMPLE]
print("This is a", labels[sample], sample)
print(X_train[MY_SAMPLE])
plt.imshow(X_train[MY_SAMPLE])
plt.show()
```

3. FASHION-ITEMS.PY (5/12)

```
# data scaling into [0, 1]
X_train = X_train / 255.0
X_test = X_test / 255.0
```

3. FASHION-ITEMS.PY (6/12)

```
# reshaping before entering CNN
# one-hot encoding is used for the output
# we use channel-last ordering (keras default)

X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
Y_train = np_utils.to_categorical(Y_train, TOTAL_CLASS)
Y_test = np_utils.to_categorical(Y_test, TOTAL_CLASS)
```

NP_UTILS.TO_CATEGORICAL()

```
from keras.utils import np_utils  
import numpy as np  
  
my_array = [0, 2, 1, 2, 0]  
  
# one-hot encoding  
one_hot = np_utils.to_categorical(my_array, 3)  
print(one_hot)
```

3. FASHION-ITEMS.PY (7/12)

```
#####
# MODEL BUILDING AND TRAINING #
#####

# build a keras sequential model of our CNN
# total parameter count formula =
# (filter_height * filter_width * input_channels + 1) * #_filters
model = Sequential()
model.add(InputLayer(input_shape = (28, 28, 1)))
```

3. FASHION-ITEMS.PY (8/12)

```
# no. parameters = (2 x 2 x 1 + 1) x 32
model.add(Conv2D(32, kernel_size = (2, 2), padding = 'same',
, activation = 'relu'))

model.add(MaxPool2D(padding = 'same', pool_size = (2, 2)))

# no. parameters = (2 x 2 x 32 + 1) x 64
model.add(Conv2D(64, kernel_size = (2, 2), padding = 'same',
, activation = 'relu'))

model.add(MaxPool2D(padding = 'same', pool_size = (2, 2)))
```

3. FASHION-ITEMS.PY (8/12)

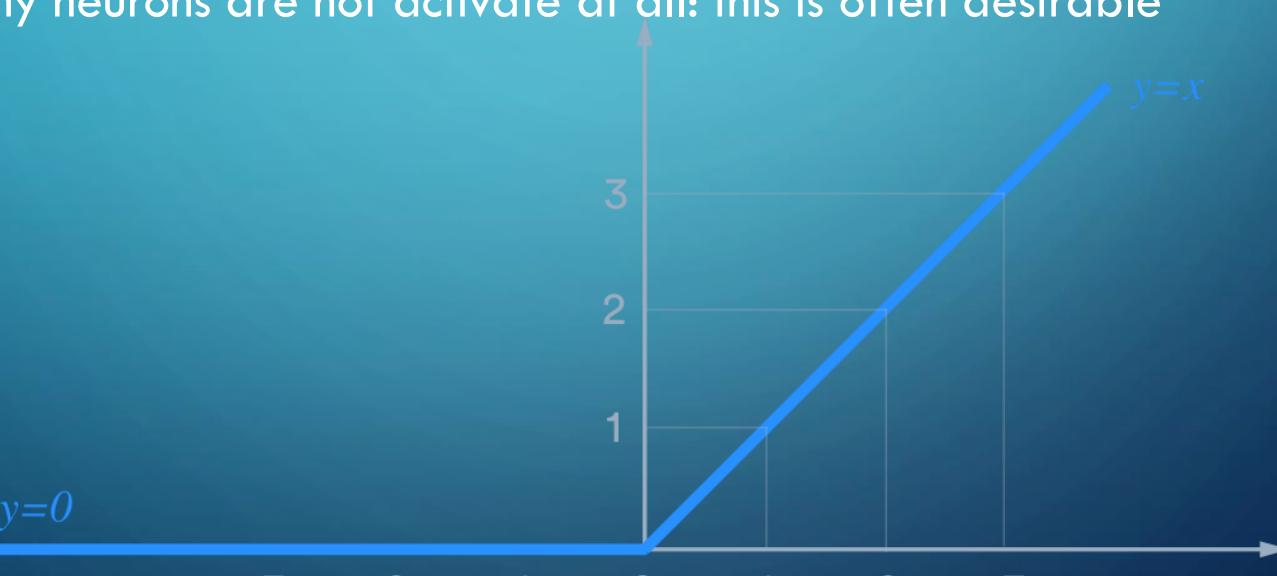
```
model.add(Flatten())
model.add(Dense(128, activation = 'relu'))

model.add(Dense(TOTAL_CLASS, activation = 'softmax'))
model.summary()
```

RELU ACTIVATION

Rectified linear unit (ReLU)

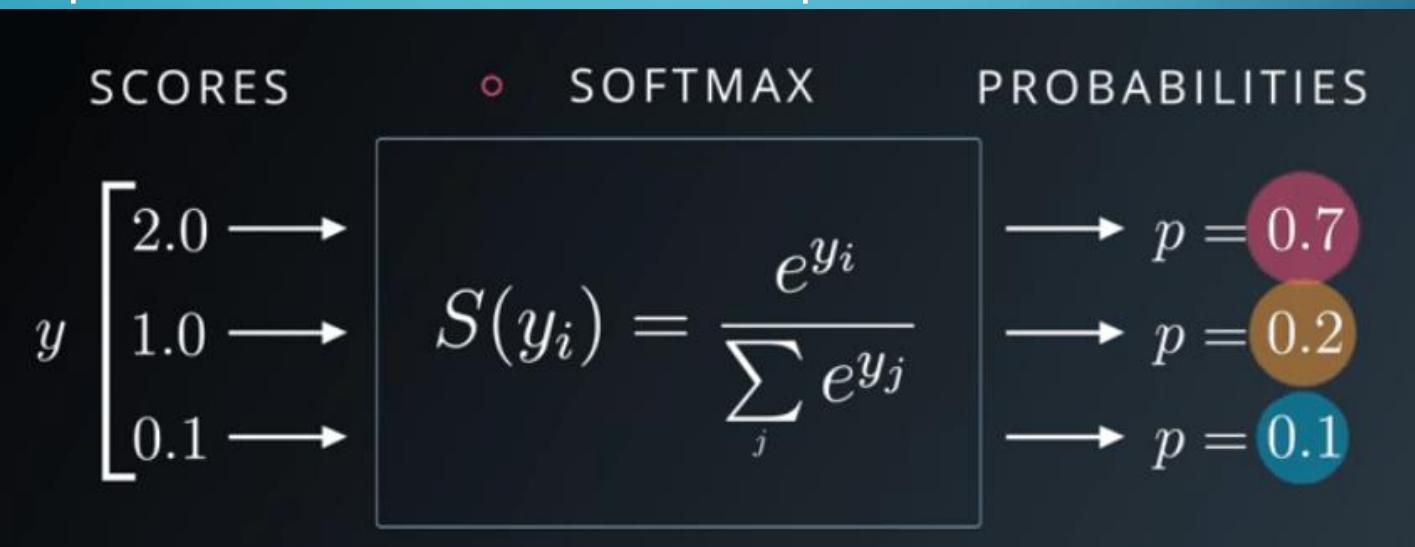
- $F(x) = \max(0, x)$
- Most commonly used activation function, especially in CNNs
- No vanishing gradient problem as in sigmoid or tanh
- Many neurons are not activate at all: this is often desirable



SOFTMAX ACTIVATION

Useful for multi-class classification

- Assigns probabilities to an object being one of several different things
- Softmax gives a list of values [0, 1] that **add up to 1**
- Step 1: we add up the evidence of our input being in certain classes
- Step 2: we convert that evidence into probabilities



3. FASHION-ITEMS.PY (9/12)

```
# model training and saving  
model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])  
model.fit(X_train, Y_train, epochs = MY_EPOCH, batch_size = MY_BATCH, verbose = 1)  
model.save('after.h5')
```

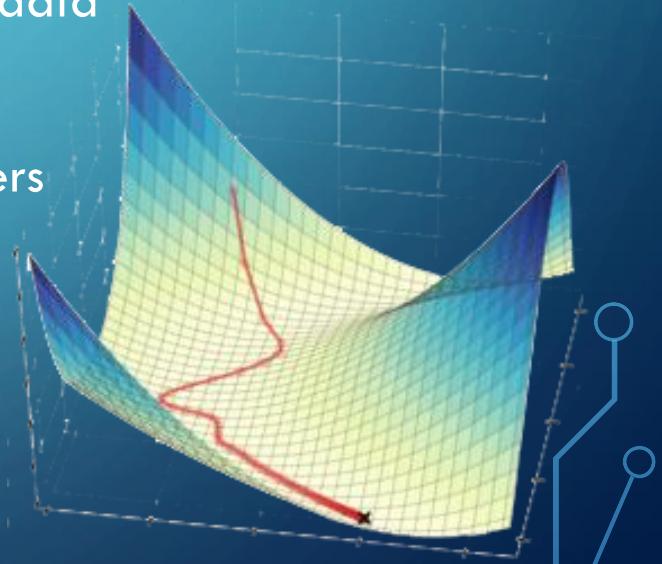
ADAM OPTIMIZER

Adam: Adaptive moment estimation

- Keeps separate learning rates for each weight
- Exponentially decay average of previous gradients
- Fairly memory efficient since it doesn't keep a history of anything
- Work well for both sparse matrices and noisy data
- Seems promising for the stock market data.

Works well with little tuning of hyperparameters

Day	Close	Change	% Change	Open	High	Low
2017-07-03	100.08	+5.97	+5.97%	94.11	100.08	93.09
2017-07-04	564.23	+2.13	+2.13%	562.10	564.23	560.00
2017-07-05	765.90	+6.43	+6.43%	759.47	765.90	760.00
2017-07-06	120.34	-11.67	-11.67%	131.00	120.34	118.00
2017-07-07	893.23	+23.11	+23.11%	870.12	893.23	860.00
2017-07-09	128.98	+5.56	+5.56%	123.42	128.98	125.00
2017-07-10	432.12	-3.67	-3.67%	435.79	432.12	425.00
2017-07-11	765.23	+11.33	+11.33%	753.90	765.23	750.00



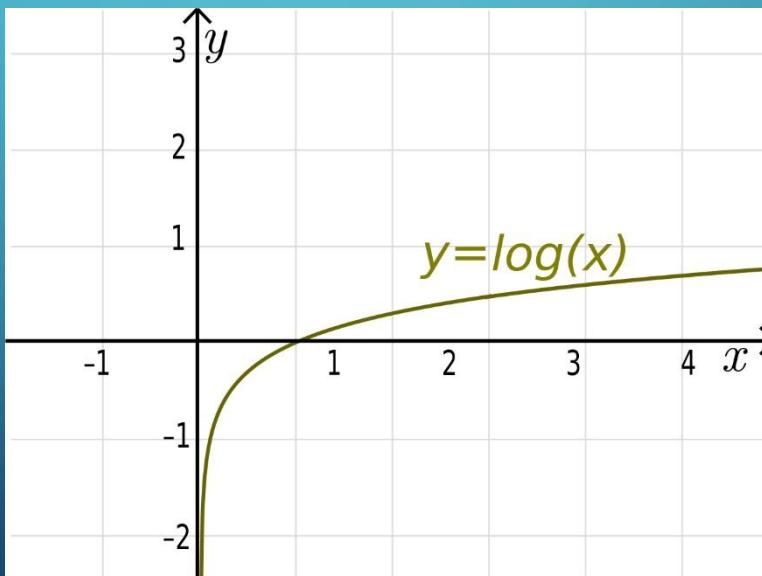
<http://ruder.io/optimizing-gradient-descent/>

CROSS ENTROPY LOSS FUNCTION

Cross entropy function

$$H_{y'}(y) = - \sum_i y'_i \times \log(y_i)$$

- y is prediction (softmax: $[0, 1]$), and y' is true value (one-hot: 0/1)
- measures how inefficient our predictions are for describing the truth



CROSENTROPY

```
# illustration of cross entropy loss function
import numpy as np

j = [0.03, 0.03, 0.01, 0.9, 0.01, 0.01, 0.0025, 0.0025,
     0.0025, 0.0025]
k = [0,0,0,1,0,0,0,0,0,0]

# cross entropy calculation
log = -(np.log(j))
print(log)
prod = k * log
print(prod)
i = np.sum(prod)
print(i)
```

3. FASHION-ITEMS.PY (10/12)

```
#####
# MODEL EVALUATION #
#####

# model evaluation
score = model.evaluate(X_test, Y_test, verbose = 1)
print('\nKeras CNN model loss = ', score[0])
print('Keras CNN model accuracy = ', score[1])
```

3. FASHION-ITEMS.PY (11/12)

```
# display confusion matrix
# we take argmax on one-hot encoding results
from sklearn.metrics import confusion_matrix
answer = np.argmax(Y_test, axis = 1)

pred = model.predict(X_test)
pred = np.argmax(pred, axis = 1)

print('\n== CONFUSION MATRIX ==')
print(confusion_matrix(answer, pred))
```

NUMPY ARGMAX()

```
import numpy as np

my_array = [[10, 14, 12], [13, 11, 15]]

# flatten the array
print(np.argmax(my_array))

# axis = 0 means along the column
print(np.argmax(my_array, axis = 0))

# axis = 1 means along the row
print(np.argmax(my_array, axis = 1))
```

CONFUSION MATRIX

Useful in multi-class classification

Making sense of the confusion matrix

	Predicted: NO	Predicted: YES
Actual: NO	50	10
Actual: YES	5	100



3. FASHION-ITEMS.PY (12/12)

```
# calculate F1 score using confusion matrix
from sklearn.metrics import f1_score
print("\nF1 score:", f1_score(answer, pred, average = 'micro'))

# testing the model after learning with a sample
sample = X_train[MY_SAMPLE]
sample = sample.reshape(1, 28, 28, 1)
pred = model.predict(sample)
print("\nPrediction after learning:",
      labels[np.argmax(pred)], "\n")
```

F1 SCORE

Popular metric for classification accuracy calculation

$$\text{precision} = \frac{\text{true positive}}{\text{true positive} + \text{false positive}}$$

$$\text{recall} = \frac{\text{true positive}}{\text{true positive} + \text{false negative}}$$

$$F1 \text{ score} = \frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$